



Sensitivity to target relaxation: the same saved training runs

(a) Relative leaderboard score



(b) Workloads reaching the target



- Schedule-Free AdamW v2 (PyTorch)
- AdEMAMix (AdamW-equiv.) (PyTorch)
- Schedule-Free AdamW (PyTorch)
- NAdamW (Baseline AlgoPerf v0.5) (JAX)
- Schedule-Free AdamW v2 (JAX)
- Muon (PyTorch)
- Schedule-Free AdamW (JAX)
- NAdamW (JAX)
- NAdamW (Tuned for ResNet) (JAX)
- Cautious NAdamW (JAX)
- Muon (JAX)
- Lion (PyTorch)
- AdEMAMix (PyTorch)
- Single Worker DiLoCo (JAX)

Lines connect sampled relaxation levels. Fixed 14-submission pool and 9 workloads; fastest-run references recomputed at each level.